Neural Network

Polese Carolina

Robotic Engineering - University of Genoa

s4862903@studenti.unige.it

Abstract—**Artificial neural networks are a fundamental approach in machine learning due to their efficiency to model complex relationships and learn meaningful representations from data. Auto-encoders, a specific class of neural networks, are widely used for dimensionality reduction, data compression, and unsupervised learning tasks. This report explores the fundamental principles of neural networks and auto-encoders, analysing their architecture, training process, and practical applications.**

Index Terms—**Artificial neural networks, auto-encoder**

I. INTRODUCTION

Artificial neural networks in the field of machine learning represent an innovative approach to tackling complex problems across various domains, from image processing to sequential data prediction. Auto-encoders are among the earliest neural network architectures designed to learn latent representations of data with an unsupervised approach. They operate by compressing data into a reduced-dimensional space and reconstructing the original input, aiming to capture relevant features and redundancies.

The importance of auto-encoders lies in their versatility: they have been successfully applied in tasks such as dimensionality reduction [2], anomaly detection [1], and deep network pre-training. Furthermore, recent advancements, such as the introduction of variational auto-encoders [3], have expanded their applications to include synthetic data generation and probabilistic learning. This report aims to provide an analysis of the fundamental concepts related to neural networks and auto-encoders, illustrating practical examples and experiments that demonstrate their effectiveness in real-world applications.

II. MATERIAL AND METHODS

A. Data processing

For this analysis, the MNIST dataset (Modified National Institute of Standards and Technology) was used, one of the most well-known datasets for image classification tasks. This dataset contains images of handwritten digits (0 to 9), where each image is gray-scale, has dimensions of 28×28 pixels, and each pixel has a value between 0 and 255, representing the intensity of the color (0 means black, 255 means white).

After loading the data, it was filtered to include only the digits 1 and 8. This reduced the dataset to two classes of interest. The filtered data were then split into distinct datasets:

- (x_{train}, y_{train}): 80% of the filtered data and contains the training images. It represent as 28×28 matrices that describe the pixel values, and the corresponding training labels (1 and 8), which represent the digits associated with the images.
- (x_{train}, y_{train}): 20% of the filtered data contains the images for the test. It represent as 28×28 matrices that describe the pixel values, and the corresponding labels for the test (1 and 8), which represent the digits associated with the images.

Then, the following steps were done to make the data compatible with the model architecture.

- *Data Normalization*: to improve performance and ensure compatibility with the model architecture, the data were normalized, converting the pixel values from their original range (0-255) to a range between 0 and 1. This transformation is essential to accelerate convergence during training and reduce the risk of numerical instability.
- *Data Reshaping*: since the MNIST images, when loaded, have the shape like a three-dimensional array ($num samples$, 28, 28), where: $num samples$ is the number of images in the dataset, and $28x28$ are the dimensions of each image. Many deep learning models, require the data to be represented in a four-dimensional format: ($num samples$, $height$, $width$, $channels$). In the case of MNIST images, the channels value is 1 because the images are greyscale. Thus, the data were reshaped from the form ($num samples$, 28, 28) to ($num samples$, 28, 28, 1) using the *reshape* command. This step explicitly adds the required depth channel, making the data compatible with the auto-encoder architecture.

B. Auto-encoder

The autoencoder is an artificial neural network model designed to learn a compact and meaningful representation of input data. The model consists of two main parts: the encoder, which receives the input data and compresses it into a lower-dimensional representation, and the decoder, which receives the compressed representation and reconstructs it into the original output data.

The autoencoder model used in this work is composed of the following parts:

- *Encoder*: the neural network that receives the input data and compresses it into a lower-dimensional representation. The encoder consists of a series of neural network layers, each applying a nonlinear transformation to the input data.
- *Decoder*: the neural network that receives the compressed representation and reconstructs it into the original output data. The decoder consists of a series of neural network layers, each applying a nonlinear transformation to the input data.
- *Autoencoder*: the complete model that combines the encoder and decoder. The autoencoder is trained to minimize the reconstruction error, which is the difference between the original input data and the reconstructed output data.

The autoencoder model operates as follows: the input data is fed to the encoder, which compresses the data into a lower-dimensional representation. The decoder receives this compressed representation and reconstructs it into the original output data. The reconstruction error is calculated as the difference between the original input data and the reconstructed output data. Finally, the model is trained to minimize the reconstruction error using an the Adam optimizer algorithm (short for "Adaptive Moment Estimation") to minimize the loss function during the training of neural networks.

C. Model evaluation

The methods for evaluating the correct implementation of the autoencoder model are different. The most intuitive method is visual inspection: it involves checking whether the digital reconstruction accurately reflects the original input data (see Fig. 1)

An alternative approach is the use of the VAF (Variance Accounted For), a specific metric for assessing the performance of autoencoders. The VAF indicates the quality of the data reconstruction by evaluating how representative and useful the compressed data is for reconstructing the original input. The VAF is computed as:

$$VAF = 1 - \left(\frac{var(\mathrm{x}_{test} - \mathrm{x}_{rec})}{var(\mathrm{x}_{test})} \right) \cdot 100$$

where: x_{test} represents the original values; x_{rec} represents the values reconstructed by the autoencoder. The calculated values for each reconstructed data point are reported in Fig. 2.

Fig. 1. Reconstruction of the digits data 1 and 8. In the first row there are the original data digit and in the second row the reconstructed data


Fig. 2. The figure represented the VAF computed for the digits values 1 and 8

Finally, another useful method for evaluating the model's implementation is to compute the information loss during reconstruction. This aspect is graphically represented in Fig. 3.

III. RESULTS AND CONCLUSION

The proper functioning of the autoencoder is demonstrated in Fig. 1, which highlights the model's ability to reconstruct the original data. This result showcases the effectiveness of the architecture and the training process in capturing the essential features of the input data while minimizing reconstruction error.

Additionally, the progressive reduction in information loss as the number of epochs increases, highlighting the steady improvement in model performance, see Fig. 3. This behavior reflects the optimization of the autoencoder's weights and biases over the time, allowing for better alignment between the encoded latent representation and the original data structure.

The analysis of the Variance Accounted For (VAF), shown in Fig. 2, provides further insights into the autoencoder's performance across different digital values. For the digital value 1, the VAF reaches 67.85%, indicating the model's good ability to capture and reconstruct this value with minimal information loss. Conversely, for the digital


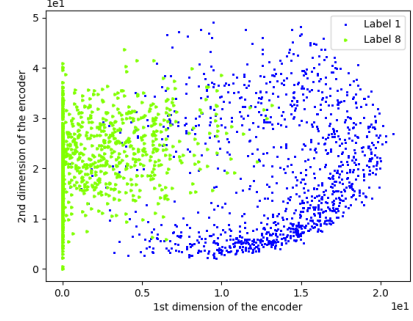Fig. 3. The figure represented the information lost during the process

Fig. 4. Visualization of the encoded latent space, showing the dimensions of the encoded representation. Points are color-coded based on digital values: Label 1 (blue) and Label 8 (green).

value 8, the VAF is 49.30%, reflecting a major difficulty in preserving information during reconstruction. These results emphasize the superior performance of the autoencoder with certain values compared to others, likely due to the complexity or distribution of the data associated with each value.

Furthermore, the scatter plot in Fig. 4 represents the latent space encoded by the autoencoder, showing the first two dimensions of the encoded representation. The points are color-coded based on the digital values 1 (blue) and 8 (green). This visualization highlights the distinct separation of the two labels in the encoded space, demonstrating the autoencoder's ability to learn meaningful representations.

The cluster corresponding to 1 appears well-defined and compact, suggesting that the autoencoder effectively captures the structure of this value. In contrast, the points associated with 8 exhibit a wider dispersion, reflecting greater variability in the encoded representation. This aligns with the previously observed lower VAF value for 8, indicating that the reconstruction process struggles more with this class.

These differences in the encoded space emphasize the variability in the autoencoder's performance depending on the complexity and distribution of the input data. Future improvements could include enhancing the encoder's ability to capture finer details for more challenging labels, potentially through increased latent dimensionality or advanced regularization techniques.

IV. DEEP LEARNING TOOLBOX

The Deep Learning Toolbox in MATLAB provides advanced tools for designing, training, and analyzing neural networks, facilitating the resolution of classification problems and other machine learning applications (see Fig. 5). Through its integrated functionalities, it is possible to visualize and evaluate model performance using confusion matrices and ROC curves. Confusion matrices offer an intuitive method to compare the predicted classes of a model with the actual ones, allowing for the analysis of classification errors and correctness. Receiver Operating Characteristic (ROC) curves, on the other hand, allow for the examination of the trade-off between the model's sensitivity and specificity, providing a graphical representation of performance across various scenarios. Using these tools, the Deep Learning Toolbox simplifies the analysis and improvement of implemented models, enabling optimization of design choices. In this study, three different neural networks were implemented:

- Neural network 1 is characterized by 2 layers: the first consisting of 10 units, and the second consisting of 2 units.
- Neural network 2 is characterized by 2 layers: the first consisting of 20 units, and the second consisting of 2 units.

Fig. 5. Graphical representation of Neural Network 2 and 3

Fig. 6. Confusion matrices for Neural Network 3

- Neural network 3 is characterized by 4 layers: the first consisting of 20 units, the second consisting of 10 units, the third consisting of 5 units, and the fourth consisting of 2 units.

To analyze the performance of each neural network, the different graphs generated by the confusion matrices and ROC curves were compared. The confusion matrices, shown in Figs. 7 and 6, highlight an improvement in model accuracy proportional to the number of layers used to define it. Examining the total confusion matrix (bottom right) for both models, it is evident that as the number of layers increases, so does the probability of correct predictions. Specifically, neural network 1 achieves an accuracy of 94.9%, while neural network 3 reaches 96.7%. This improvement suggests that more complex architectures, if well-designed, can lead to better predictive performance. The ROC graphs (see Figs. 8 and 9) confirm these findings: in neural network 3, the combined curve (bottom right) is closer to the upper-left corner, representing the ideal value, compared to that of neural network 2. Overall, the models defined in our analysis demonstrate good predictive capabilities.

REFERENCES

[1] R. Chalapathy and S. Chawla, "Deep learning for anomaly detection: A survey," *ArXiv*, vol. abs/1901.03407, 2019. [Online]. Available:
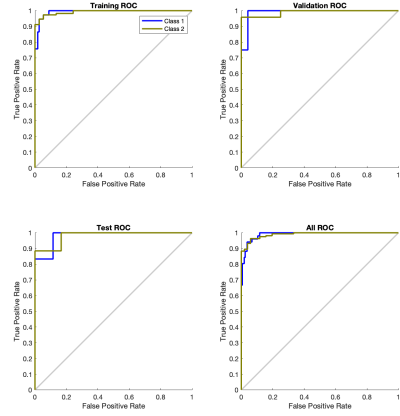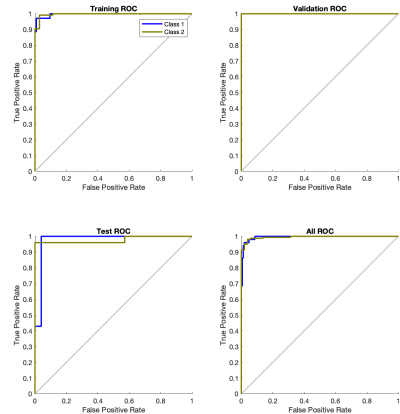

Fig. 7. Confusion matrices for Neural Network 1

Fig. 8. ROC curve for Neural Network 2

Fig. 9. ROC curve for Neural Network 3

https://api.semanticscholar.org/CorpusID:57825713

[2] G. E. Hinton and R. R. Salakhutdinov, "Reducing the dimensionality of data with neural networks," *Science*, vol. 313, no. 5786, pp. 504–507, 2006. [Online]. Available: https://www.science.org/doi/abs/10.1126/science.1127647

[3] D. P. Kingma and M. Welling, "Auto-encoding variational bayes," *CoRR*, vol. abs/1312.6114, 2013. [Online]. Available: https://api.semanticscholar.org/CorpusID:216078090